

Team 5 Grades and Feedback

Team title

Team 5: PersonaPath: Personalized Travel & Dining Recommendation Engine (Behavioral Profiling)

Members

Luong,Dat Vy;Tang,Yiwei;Zhang,Jiaqing;Konda,Amulya;Gaikwad,Pranjali Ashok

Comment AI summary

Positives

- Polished, professional UI and presentation materials, including a clear flyer/poster and well-organized GitHub, that make the product approachable.
- Innovative multi-layer recommendation engine that combines user and restaurant profiles, query intent, and a quality score to produce tailored results.
- Advanced technical elements such as entropy-adjusted sentiment, JSD vector similarity, and an intelligence layer that elevate the analysis beyond basic approaches.
- Strong personalization features—persona classification, LLM integration, feedback loop, and context-aware recommendations—that adapt as users provide input.
- Clear product focus and USP demonstrated through direct comparisons to existing recommendation sites and a compelling demo.
- User-facing features like a dynamic dashboard and a "hidden gem" capability that surface relevant, less-known options.

Major suggestions

- Improve onboarding and first-time user experience (e.g., provide a working quick-start guide and consider an LLM-powered chat to capture preferences).
- Replace or augment LDA with modern embedding-based models and benchmark multiple modeling approaches to better represent user/restaurant "DNA."
- Add guardrails and fallback behavior for missing or sparse data and strengthen validation of recommendation scoring by comparing alternative algorithms beyond cosine similarity.
- Clarify and document how different preference signals are weighted and expose controls or explanations so users understand and customize recommendation behavior.
- Expand geographic coverage and real-world testing beyond a single default city, including GPS/location integration and additional city datasets.
- Prioritize productization and go-to-market considerations to move the prototype toward deployment and broader adoption.

Grading details

Dimension	Grade or %
Peer - Business value (out of 5)	4.5
Peer - Technical quality (out of 5)	4.8
Peer - Effectiveness of presentation (out of 5)	4.8

Dimension	Grade or %
Peer - Professionalism (out of 5)	4.9
Total Peer (out of 100)	95.2
Peer - bestinshow (% of evaluators voted for this)	62.5%
Instructor - Business value (out of 10)	9.5
Instructor - Technical quality (out of 10)	9.5
Instructor - Effectiveness of presentation (out of 10)	9.5
Instructor - Creativity (out of 10)	9.5
Total Instructor (out of 100)	95
Judge - Business value (out of 5)	4.4
Judge - Technical quality (out of 5)	4.2
Judge - Ppresentation (out of 5)	4.2
Judge - Handout (out of 5)	4.4
Judge - Professionalism (out of 5)	4.6
Judge - Best in Show (%)	60%
Overall - Best in Show (%)	61%
Judge Total (out of 100)	86.6
Deliverable Completion (out of 100)	100
Combined Rating (out of 100)	93.5
Combined Rating (out of 100) w/ Best in show bonus	93.5
Team Grade	14.0

Instructor comments

Business value

- personalized assistance/recommendations are important

Technical Quality

- traditional dataset but innovative analysis (intelligence layer, enhanced sentiment, JSD vector similarity search) - streamlit UI

Presentation

- compelling and concise flyer - well organized github

Novelty

- innovative 4-layer engine (that combined restaurant/person DNA, query intent, and Quality score) - use Entropy-Adjusted Sentiment

Overall evaluation (Verdict)

- - novel combination of newer techniques to generate new/improved solutions to a traditional problem

Weakness / Area of improvements

- - LDA may not be a strong technique for modeling DNA. modern embedding tend to work better
- quick start guide missing (broken link)

Judge comments

I like

- The idea of using ones own reviews to build a profile. The UI was really professional.
- Identified a real problem
- Presentation was compelling. I want one.
- I like the persona effect in this application
- I like how the team presented a clean and intuitive UI for the tool. I also like how they compared it with other restaurant recommendation websites and clearly highlighted their USP, which helped position their product effectively.

I wish

- More strong way to on board new users to the system.
- Instead of preloaded limited user profiles they could provide an llm powered chat experience that is open to capture user preferences.
- Start selling it.
- I wish they had added guardrails to prevent the tool from showing results when data is missing for a user query. I also wish they had explored more approaches to make the recommendation system more robust, such as comparing traditional methods before moving to cosine scores. Additionally, I wish the team had been more coherent while answering questions during the demo.

Peer comments

I like

- They have a feedback system to improve their performance.
- Persona classification
- I like how valuable the problem they're solving. It's very impressive
- I like how they integrated the recommendations system with their llm
- I liked that this was a tool I could see myself using if I'm unsure where I want to go to eat. This seems like a very relevant and useful tool for myself.
- I like how the dashboard keeps updating with data you input. If you dislike a restaurant that was recommended, it is less likely to recommend that to you
- I like how the recommendations are based on context-aware personalization, which is often more aligned with how people actually search.
- I like there poster design: very aesthetic, clear information.
- that restaurant recommendations are personalized to the user and that you could make requests in plain language

- As someone who enjoys traveling and finding new spots to eat, I would find this extremely helpful.
- They did a great job with their presentation and showing the direct value of who it is meant for
- Their handout was very easy to follow. They walked through their process easily.
- The idea was fantastic and the execution as well
- The presentation and the front end of the project was really polished.
- The team's UI looks very appealing and easy to use. The team did a great job explaining the technological pipeline.
- There is a fun feature called hidden gem, so it actually put the good local restaurants under the spotlight.

I wish

- It is only tested in a default town, not in a gps based on my location
- ml model could have been better
- I wish they'd turned this into a reality and build the app for it.
- Generalized it to national map, not only Philadelphia
- I wish it was done for the Minneapolis area instead of Philadelphia so I could actually get really recommendations. If the data wasn't available though, then that makes sense.
- I would love to see more cities shown
- I think they could further differentiate their approach from existing platforms like Google Maps, since users can already search for contextual keywords and receive recommendations based on review content.
- I wish they would be clearer in how their algorithm incorporate different sources of users' pref into generating recommendations (how much weight they put into each sources users put in - preference, rating in app, etc.)
- there was a bit more innovation/differentiation. This product seems quite similar to existing restaurant search engines.
- I think they could improve upon the user rating system. If they vote thumbs up or thumbs down, perhaps an additional area to say why they rated a specific restaurant that way to feed into the LLM.
- I think it would've been awesome to have a sample size of some different cities, but they promptly addressed that!
- I wish they could extend it to beyond yelp. I think it would be cool to incorporate item recommendations as well
- they could have given more customisation option
- I feel like there are things out there that can do this, maybe not as personalized, but still available in some degree.
- There are apps and websites that solve address the same problem so the solution didn't stand out as unique.
- Scale up the geographical location.